

Part 2: iOS Deploy



Select files to import

To build a custom model, you first need to import a set of images to train it. Each image should be categorized with a label. (Labels are essential for telling the model how to identify an image.)

- Each label should have at least 100 images for best results.

☒ Upload images from your computer

☐ Select a CSV file on Cloud Storage

Upload images from your computer

Supports JPG, PNG, GIF, BMP, ICO, ZIP. Maximum 500 files per upload. Uploaded files will be stored on Cloud Storage.

flower_photos.zip

1 file



SELECT FILES

Destination on Cloud Storage



gs:// cloud-ai-platform-91507f72-3075-45fc-8750-ff5c73f5aa44/flow

BROWSE

CONTINUE



AI Platform (Unified)

test-dataset

test-dataset_1cn



TRAIN NEW MODEL

CREATE LABELING TASK



Dashboard



Datasets



Labeling tasks



Notebooks



Training



Models



Endpoints



Batch predictions

IMPORT

BROWSE

ANALYZE

All 3,667

Labeled 3,667

Unlabeled 0

Filter labels



daisy 633

dandelion 898

roses 641

sunflowers 697

tulips 798

ADD NEW LABEL

Filter items



Unable to import data due to errors.

DETAILS

DISMISS

☐ Select all

roses



daisy



tulips



daisy



sunflowers





AI Platform (Unified)



test-dataset_20212160186

VIEW DATASET



Dashboard



Datasets



Labeling tasks



Notebooks



Training



Models



Endpoints



Batch predictions

EVALUATE

DEPLOY & TEST

BATCH PREDICTIONS

MODEL PROPERTIES

Filter labels



All labels	0
dandelion	0.99975
sunflowers	0.99901
daisy	0.99721
tulips	0.99520
roses	0.99126

Confidence threshold 0.5

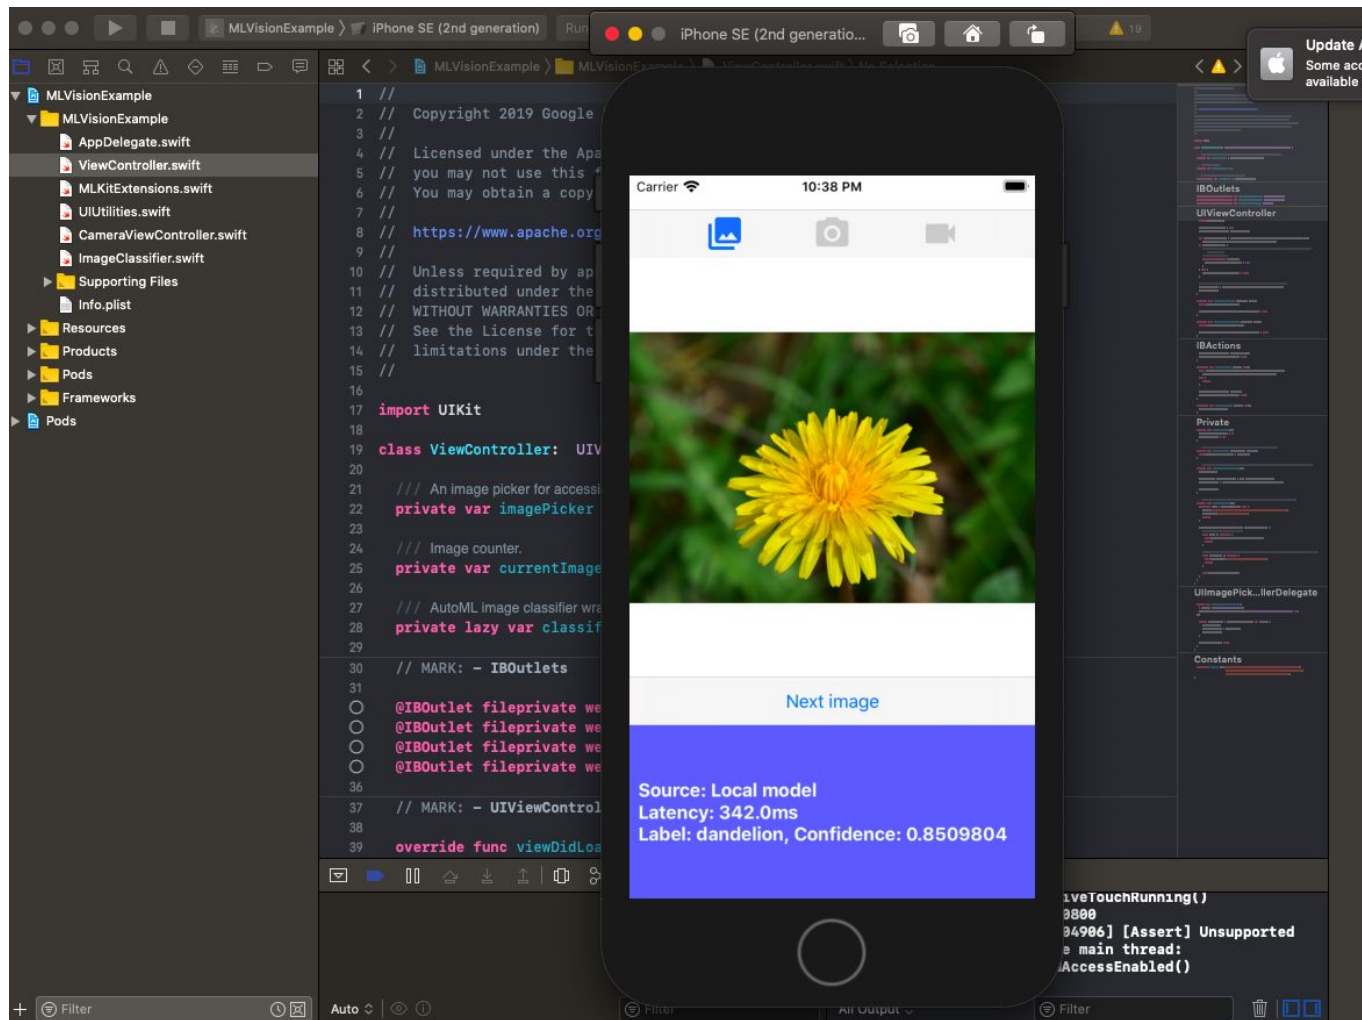
All labels

Average precision	0.996
Precision	97.3%
Recall	97%
Created	Feb 15, 2021, 4:44:52 PM
Total images	3,667
Training images	2,933
Validation images	367
Test images	367

Use the slider to see which score threshold works best for your model on the precision-recall tradeoff curve. [Learn more about these metrics and graphs](#)



```
mlkit-automl — shannon@Shannons-MacBook-Air — zsh — 80x36
../mlkit-automl      ../mlkit-automl  +
mlkit-automl % pod install --verbose --no-repo-update
Preparing
Analyzing dependencies
Inspecting targets to integrate
Using 'ARCHS' setting to build architectures of target 'Pods-MLVisionExample':
(``)
Finding Podfile changes
- FirebaseMLCommon
- FirebaseMLVision
- FirebaseMLVisionAutoML
Resolving dependencies of `Podfile`
No Editor
No Selection
Comparing resolved specification to the sandbox manifest
A FirebaseCore
A FirebaseInstanceID
A FirebaseMLCommon
A FirebaseMLVision
A FirebaseMLVisionAutoML
A GTMSessionFetcher
A GoogleAPIClientForREST
A GoogleMobileVision
A GoogleToolboxForMac
A GoogleUtilities
A Protobuf
A TensorFlowLite
Downloading dependencies
-> Installing FirebaseCore (6.0.0)
> Git download
> Git download
$ /usr/local/bin/git clone https://github.com/firebase/firebase-ios-sdk.git
```



Part 2: Time Series

Google Cloud Platform

My First Project

Search products and resources

AI Platform

Dashboard

AI Hub

Data Labeling

Notebooks

Pipelines

Jobs

Models

Notebooks

Create and use Jupyter Notebooks with a notebook instance. Notebook instances have JupyterLab pre-installed and are configured with GPU-enabled machine learning frameworks. [Learn more](#)

Filter table

	Instance name	Zone	Environment	Machine type	GPUs	Permission	Labels
<input type="checkbox"/>	tensorflow-2-1-20210219-200805	us-west1-b	TensorFlow:2.1	4 vCPUs, 15 GB RAM	None	Service account	No labels

Setting up proxy to JupyterLab

Info panel

DOCUMENTATION

LABELS

[Notebook instances](#)

[Notebook API](#)

7765e84c98dfc01c-dot-us-west1.notebooks.googleusercontent.com/lab?authuser=1

Apps

searchstc-svc-c09 [...]

SVMT Menu

Scope

Satori

U3

Tansa

日本語

Commute Calendar...

OneDrive

Masters

MSR AI Seminars [...]

EY Eugene Yan

EXP Related Categories.

File

Edit

View

Run

Kernel

Git

Tabs

Settings

Help

AI Platform Notebooks

/

Name

Last Modified

src

3 minutes ago

training-data-analyst

seconds ago

tutorials

3 minutes ago

Untitled.ipynb

seconds ago

Untitled.ipynb

Terminal 1

(base) jupyter@tensorflow-2-1-20210219-200805:~\$ git clone https://github.com/GoogleCloudPlatform/training-data-analyst
Cloning into 'training-data-analyst'...
remote: Enumerating objects: 45287, done.
remote: Total 45287 (delta 0), reused 0 (delta 0), pack-reused 45287
Receiving objects: 100% (45287/45287), 477.96 MiB | 31.51 MiB/s, done.
Resolving deltas: 100% (28483/28483), done.
Checking out files: 100% (9153/9153), done.
(base) jupyter@tensorflow-2-1-20210219-200805:~\$

Untitled.ipynb × Terminal 1 × 01-explore.ipynb


- Is there a difference in ridership between the weekday and weekends?
- Is the mix of bus vs rail ridership changing over time?

```
[15]: # Initialize plotting
```

```
register_matplotlib_converters() # Addresses a warning
sns.set(rc={'figure.figsize':(16,4)})
```

```
[16]: # Explore total rides over time
```

```
sns.lineplot(data=df, x=df.index, y=df[target]).set_title('Total Rides')
fig = plt.show()
```



```
[17]: # Explore rides by day type: Weekday (W), Saturday (A), Sunday/Holiday (U)
```

```
sns.lineplot(data=df, x=df.index, y=df[target], hue=df['day_type']).set_title('Total Rides by Day Type')
fig = plt.show()
```



```
[18]: # Explore rides by transportation type
```

Google Cloud Platform

My First Project

Search products and resources

BigQuery

SQL workspace

Data transfers

Scheduled queries

Reservations

BI Engine

FEATURES & INFO

SHORTCUT

HIDE PREVIEW FEATURES

Explorer

+ ADD DATA

Type to search

Viewing pinned projects.

round-dreamer-304721

demo

MORE RESULTS

EDITOR 2

ROUND-...

DEMO

DEMO

COMPOSE NEW QUERY

round-dreamer-304721:demo

CREATE TABLE

SHARE DATASET

AUTHORIZE ROUTINES

COPY DATASET

DELETE DATASET

Description

None

Labels

None

Dataset info

Dataset ID

Created

Default table expiration

Last modified

Data location

round-dreamer-304721:demo

Feb 19, 2021, 8:16:49 PM

5 days 0 hr

Feb 19, 2021, 8:16:49 PM

US

JOB HISTORY

QUERY HISTORY

SAVED QUERIES

Job history

REFRESH

Personal history

Project history

Filter queries

Today

8:17 PM

Load uploaded file to round-dreamer-304721:demo.cta_ridership

Google Cloud Platform

My First Project

Search products and resources

BigQuery

SQL workspace

Data transfers

Scheduled queries

Reservations

BI Engine

FEATURES & INFO

SHORTCUT

HIDE PREVIEW FEATURES

Explorer

+ ADD DATA

Type to search

Viewing pinned projects.

round-dreamer-304721

demo

cta_ridership

cta_ridership_model

MORE RESULTS

MORE RESULTS

EDITOR 2

ROUND...

DEMO

DEMO

CTA_RID...

UNSAVE... 3

CTA_RID...

CTA_RID...

COMPOSE NEW QUERY

RUN

SAVE

SCHEDULE

MORE

This query will process 108.4 KiB (ML) when run.

```
1 CREATE OR REPLACE MODEL
2   `demo.cta_ridership_model` OPTIONS(MODEL_TYPE='ARIMA',
3   `TIME_SERIES_TIMESTAMP_COL`='service_date',
4   `TIME_SERIES_DATA_COL`='total_rides',
5   `HOLIDAY_REGION`='us') AS
6 SELECT
7   service_date, total_rides
8 FROM
9   `demo.cta_ridership`
```

Query results

Query complete (39.9 sec elapsed, 4.4 MB (ML) processed)

Job information

Results

Execution details

This statement will create a new model named round-dreamer-304721.demo.cta_ridership_model. Depending on the type of model, this may take several hours to complete.

Go to model

BigQuery

SQL workspace

Data transfers

Scheduled queries

Reservations

BI Engine

FEATURES & INFO

SHORTCUT

HIDE PREVIEW FEATURES

Explorer

+ ADD DATA

Type to search

Viewing pinned projects.

round-dreamer-304721

demo

cta_ridership

cta_ridership_model

MORE RESULTS

MORE RESULTS

cta_ridership_model

QUERY MODEL

DELETE MODEL

EXPORT MODEL

DETAILS

TRAINING

EVALUATION

SCHEMA

Non Seasonal P	Non Seasonal D	Non Seasonal Q	Has Drift	Log Likelihood	AIC	Variance	Seasonal Period
1	1	4	True	-84,343.913	168,701.826	2,121,476,632.467	Weekly, Yearly
1	1	4	False	-84,345.763	168,703.526	2,122,628,259.179	Weekly, Yearly
4	1	1	True	-84,346.869	168,707.738	2,123,285,308.131	Weekly, Yearly
1	1	3	True	-84,347.973	168,707.946	2,123,959,900.714	Weekly, Yearly
4	1	1	False	-84,348.833	168,709.666	2,124,511,101.972	Weekly, Yearly
1	1	3	False	-84,349.844	168,709.688	2,125,133,805.121	Weekly, Yearly
3	1	2	True	-84,349.466	168,712.932	2,124,875,662.833	Weekly, Yearly
2	1	2	True	-84,351.446	168,714.892	2,126,093,471.221	Weekly, Yearly
2	1	2	False	-84,352.879	168,715.758	2,126,992,389.241	Weekly, Yearly
0	1	5	True	-84,351.067	168,716.134	2,125,867,126.489	Weekly, Yearly
0	1	5	False	-84,352.172	168,716.343	2,126,558,119.381	Weekly, Yearly
3	1	2	False	-84,352.526	168,717.053	2,126,777,555.452	Weekly, Yearly
2	1	3	True	-84,354.875	168,723.749	2,128,191,178.209	Weekly, Yearly
2	1	3	False	-84,357.564	168,727.128	2,129,867,036.91	Weekly, Yearly
3	1	1	True	-84,366.456	168,744.913	2,135,328,985.565	Weekly, Yearly
3	1	1	False	-84,367.982	168,745.963	2,136,287,297.784	Weekly, Yearly
2	1	1	False	-84,407.403	168,822.807	2,160,766,577.274	Weekly, Yearly
2	1	1	True	-84,406.65	168,823.3	2,160,286,811.78	Weekly, Yearly
0	1	4	False	-84,416.093	168,842.186	2,166,260,904.36	Weekly, Yearly
0	1	4	True	-84,415.732	168,843.463	2,166,031,399.586	Weekly, Yearly
0	1	3	False	-84,433.324	168,874.648	2,177,098,384.353	Weekly, Yearly
0	1	3	True	-84,433.107	168,876.215	2,176,960,425.757	Weekly, Yearly

- BigQuery
- SQL workspace
- Data transfers
- Scheduled queries
- Reservations
- BI Engine

FEATURES & INFO

SHORTCUT

HIDE PREVIEW FEATURES

Explorer

+ ADD DATA

Type to search

Viewing pinned projects.

round-dreamer-304721

demo

cta_ridership

cta_ridership_model

MORE RESULTS

MORE RESULTS

ROUND-... DEMO DEMO CTA_RID... UNSAVE... 3 CTA_RID... CTA_RID... COMPOSE NEW QUERY

RUN SAVE SCHEDULE MORE

```
1 SELECT
2 | *
3 FROM
4 | ML_EVALUATE(MODEL `demo.cta_ridership_model`)
```

This query will process 0 B when run.

Query results

SAVE RESULTS EXPLORE DATA

Query complete (0.2 sec elapsed, 0 B processed)

Job information Results JSON Execution details

Row	non_seasonal_p	non_seasonal_d	non_seasonal_q	has_drift	log_likelihood	AIC	variance	seasonal_periods
1	1	1	4	true	-84343.91298029698	168701.82596059397	2.1214766324672794E9	WEEKLY
								YEARLY
2	1	1	4	false	-84345.76278035615	168703.5255607123	2.1226282591786644E9	WEEKLY
								YEARLY
3	4	1	1	true	-84346.86918283005	168707.7383656601	2.1232853081307085E9	WEEKLY
								YEARLY
4	1	1	3	true	-84347.97278479983	168707.94556959966	2.1239599007139666E9	WEEKLY
								YEARLY
5	4	1	1	false	-84348.83291975319	168709.66583950637	2.124511101972134E9	WEEKLY
								YEARLY
6	1	1	3	false	-84349.84391557962	168709.68783115924	2.1251338051213825E9	WEEKLY

Rows per page: 100 1 - 42 of 42 First page < > > Last page

anon04fe69279440cfb354852927cb6abdac03ee851

created. Go to table

- BigQuery
- SQL workspace
- Data transfers
- Scheduled queries
- Reservations
- BI Engine

FEATURES & INFO

SHORTCUT

HIDE PREVIEW FEATURES

Explorer

+ ADD DATA

Type to search

Viewing pinned projects.

round-dreamer-304721

demo

cta_ridership

cta_ridership_model

MORE RESULTS

MORE RESULTS

ROUND-... DEMO DEMO CTA_RID... *UNSAVE... 3 CTA_RID... CTA_RID... COMPOSE NEW QUERY

RUN SAVE SCHEDULE MORE

```
1 SELECT
2 | *
3 FROM
4 | ML_FORECAST(MODEL `demo.cta_ridership_model`,
5 | STRUCT(7 AS horizon))
```

Query results

SAVE RESULTS EXPLORE DATA

Query complete (0.2 sec elapsed, 23.4 KB processed)

Job information Results JSON Execution details

Row	forecast_timestamp	forecast_value	standard_error	confidence_level	prediction_interval_lower_bound	prediction_interval_upper_bound	confidence_interval_lower_bound	confidence_interval_upper_bound
1	2020-01-01 00:00:00 UTC	662436.4424369269	46059.49014554253	0.95	572322.980240453	752549.9046334007	572322.980240453	752549.9046334007
2	2020-01-02 00:00:00 UTC	1029641.4669424891	46276.328347693256	0.95	939103.76989082	1120179.1639941582	939103.76989082	1120179.1639941582
3	2020-01-03 00:00:00 UTC	1201660.2034356925	47233.43871922012	0.95	1109249.9600529654	1294070.4468184195	1109249.9600529654	1294070.4468184195
4	2020-01-04 00:00:00 UTC	651095.9776391207	48157.99332862347	0.95	556876.8819095747	745315.0733686666	556876.8819095747	745315.0733686666
5	2020-01-05 00:00:00 UTC	467394.91846646497	48621.50963880497	0.95	372268.97250121285	562520.8644317171	372268.97250121285	562520.8644317171
6	2020-01-06 00:00:00 UTC	1158999.319539823	48869.23710364581	0.95	1063388.705171438	1254609.9339082083	1063388.705171438	1254609.9339082083
7	2020-01-07 00:00:00 UTC	1127789.5651062205	49011.66149084522	0.95	1031900.3033930386	1223678.8268194026	1031900.3033930386	1223678.8268194026

File Edit View Run Kernel Git Tabs Settings Help AI Platform Notebooks

+
-
↑
↺
↻

... / ai-for-time-series / notebooks /

Name	Last Modified
cnr_export	a minute ago
data	18 minutes ago
lstm_export	2 minutes ago
01-explore.ipynb	3 minutes ago
02-model.ipynb	2 minutes ago
03-cloud-training.ipynb	6 minutes ago
cta_ridership.csv	16 minutes ago

Untitled.ipynb Terminal 1 01-explore.ipynb 02-model.ipynb 03-c

Markdown git

In this section, you will build models using popular neural network architectures for time-series data.

Long Short Term Memory (LSTM)

```
[45]: # Reshape test data to match model inputs and outputs

X_train = X_train_reframed.values.reshape(-1, n_input_steps, n_features)
X_test = X_test_reframed.values.reshape(-1, n_input_steps, n_features)
y_train = y_train_reframed.values.reshape(-1, n_output_steps, 1)
y_test = y_test_reframed.values.reshape(-1, n_output_steps, 1)
```

TODO 2: Update the LSTM architecture

Try increasing and decreasing the number of LSTM units and see if you notice any accuracy improvements.

You can use hyper-parameter tuning to search for optimal values, but that's outside the scope of this lab.

```
[46]: # Try increasing and decreasing the number of LSTM units and see if you notice any accuracy improvements.
# Run the next cell to evaluate the results in more detail.

model = Sequential([
    LSTM(64, input_shape=[n_input_steps, n_features]),
    Dense(n_output_steps)])

model.compile(optimizer='adam', loss='mae')

early_stopping = EarlyStopping(monitor='val_loss', patience=patience)
_ = model.fit(x=X_train, y=y_train, validation_data=(X_test, y_test), epochs=epochs, callbacks=[early_stopping])

Train on 5515 samples, validate on 1352 samples
Epoch 1/1000
5515/5515 [=====] - 10s 2ms/sample - loss: 0.6554 - val_loss: 0.4750
Epoch 2/1000
5515/5515 [=====] - 6s 1ms/sample - loss: 0.3387 - val_loss: 0.2568
Epoch 3/1000
5515/5515 [=====] - 7s 1ms/sample - loss: 0.2406 - val_loss: 0.2406
Epoch 4/1000
5515/5515 [=====] - 7s 1ms/sample - loss: 0.2260 - val_loss: 0.2383
Epoch 5/1000
5515/5515 [=====] - 7s 1ms/sample - loss: 0.2207 - val_loss: 0.2263
Epoch 6/1000
5515/5515 [=====] - 6s 1ms/sample - loss: 0.2156 - val_loss: 0.2218
Epoch 7/1000
5515/5515 [=====] - 6s 1ms/sample - loss: 0.2106 - val_loss: 0.2208
Epoch 8/1000
5515/5515 [=====] - 8s 1ms/sample - loss: 0.2071 - val_loss: 0.2352
Epoch 9/1000
5515/5515 [=====] - 7s 1ms/sample - loss: 0.2052 - val_loss: 0.2142
Epoch 10/1000
5515/5515 [=====] - 7s 1ms/sample - loss: 0.2025 - val_loss: 0.2181
Epoch 11/1000
5515/5515 [=====] - 7s 1ms/sample - loss: 0.2016 - val_loss: 0.2163
```

1 4 tensorflow-2-1-20210219:200805 | round-dreamer-304721 Git: refreshing... Python 3 | Busy

File Edit View Run Kernel Git Tabs Settings Help AI Platform Notebooks

+

+

+

+

+

/ ... / ai-for-time-series / notebooks /

Name	Last Modified
cnm_export	2 minutes ago
data	19 minutes ago
lstm_export	3 minutes ago
01-explore.ipynb	4 minutes ago
02-model.ipynb	seconds ago
03-cloud-training.ipynb	6 minutes ago
cta_ridership.csv	17 minutes ago

Untitled.ipynb Terminal 1 01-explore.ipynb 02-model.ipynb 03-cloud-training.ipynb

Markdown git

```
[47]: model.save('./lstm_export/')

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/resource_variable_ops.py:1786: calling Ba
constraint is deprecated and will be removed in a future version.
Instructions for updating:
If using Keras pass *_constraint arguments to layers.
INFO:tensorflow:Assets written to: ./lstm_export/assets

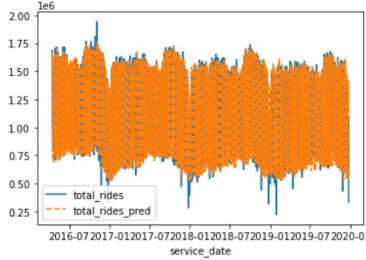
[48]: # Predict the results, and then reverse the transformation that scaled all values to a mean of 0 and std. dev. of 1
preds = model.predict(X_test)
y_pred_lstm = inverse_scale(preds)

# Evaluate the overall results and for each time step
evaluate(y_pred_lstm)

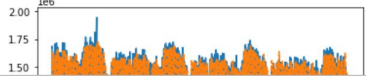
# The plot will show the R^2 value (0 lowest -> 1 highest) and the MAE (mean absolute error) for the entire prediction window.
# It will also show individual plots for 1 day out, 2 days out, etc. comparing the actual vs the predicted value.

=== t+(1-7) ===
R^2: 0.812
MAPE: 0.092
MAE: 81117.131

=== t+1 ===
R^2: 0.858
MAPE: 0.079
MAE: 69010.06
```



```
=== t+2 ===
R^2: 0.827
MAPE: 0.089
MAE: 78629.465
```



1 4 tensorflow-2-1-20210219-200805 | round-dreamer-304721 Git: refreshing... Python 3 | Busy Saving completed